# Pingzhi Li

**Email**: pingzhi@cs.unc.edu **Website**: pingzhili.github.io **Location**: NC, US

## Education

**The University of North Carolina at Chapel Hill** NC, US
Ph.D. in Computer Science Sep. 2024 – Jul. 2029 (Estimated)

**University of Science and Technology of China** Hefei, China
Bachelor of Engineering in Computer Science Sep. 2019 – Jul. 2023

## Publications

(* Equal Contribution) (ˆ Equal Supervision)

**Model-GLUE: Democratized LLM Scaling for A Large Model Zoo in the Wild**
**Pingzhi Li***, Xinyu Zhao*, Guoheng Sun*, Ruisi Cai*, Yukun Zhou*, Peihao Wang*, Bowen Tan, Yexiao He, Li Chen, Yi Liang, Beidi Chen, Binhang Yuan, Hongyi Wangˆ, Ang Liˆ, Zhangyang Wangˆ, Tianlong Chenˆ
*Conference on Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track, 2024*
[Code] [PDF]

**Revisiting Zeroth-Order Optimization for Memory-Efficient LLM Fine-Tuning: A Benchmark**
**Pingzhi Li***, Yihua Zhang*, Junyuan Hong*, Jiaxiang Li*, Yimeng Zhang, Wenqing Zheng, Pin-Yu Chen, Jason D. Lee, Wotao Yin, Mingyi Hong, Zhangyang Wang, Sijia Liu, Tianlong Chen
*International Conference on Machine Learning (ICML), 2024*
[Code] [PDF]

**Merge, Then Compress: Demystify Efficient SMoE with Hints from Its Routing Policy**
**Pingzhi Li**, Zhenyu Zhang, Prateek Yadav, Yi-Lin Sung, Yu Cheng, Mohit Bansal, Tianlong Chen
*International Conference on Learning Representations (ICLR), 2024* (Spotlight)
[Code] [PDF]

## Under Review

(* Equal Contribution)

**Glider: Global and Local Instruction-Driven Expert Router**
**Pingzhi Li***, Prateek Yadav*, Jaehong Yoon, Jie Peng, Yi-Lin Sung, Mohit Bansal, Tianlong Chen
Submitted to *International Conference on Learning Representations (ICLR), 2025*

**Examining Post-Training Quantization for Mixture-of-Experts: A Benchmark**
**Pingzhi Li***, Xiaolong Jin*, Yu Cheng, Tianlong Chen
Submitted to *International Conference on Learning Representations (ICLR), 2025*

**Hybrid Quantum-Classical Scheduling for Accelerating Neural Network Training with Newton's Gradient Descent**
**Pingzhi Li**, Junyu Liu, Hanrui Wang, Tianlong Chen
Submitted to *International Symposium on High-Performance Computer Architecture (HPCA), 2025*

**Understanding Prejudice and Fidelity of Diverge-to-Converge Multi-Agent Systems**
Zhen Tan, Song Wang, Shyam Marjit, Zihan Chen, Yinhan He, Xinyu Zhao, **Pingzhi Li**, Jundong Li, Huan Liu, Tianlong Chen
Submitted to *International Conference on Learning Representations (ICLR), 2025*

**Fantastic Experts and How to Find Them: A Multi-Dimensional Study for Experts-Level Sparsification in Mixture-of-Experts**
Ajay Kumar Jaiswal, Jianyu Wang, Yixiao Li, **Pingzhi Li**, Tianlong Chen, Zhangyang Wang, Chong Wang, Ruoming Pang, Xianzhi Du
Submitted to *International Conference on Learning Representations (ICLR), 2025*

**PortLLM: Personalizing Evolving Large Language Models with Training-Free and Portable Model Patches**
Rana Shahroz, **Pingzhi Li***, Sukwon Yun*, Zhenyu Wang, Shahriar Nirjon, Chau-Wai Wong, Tianlong Chen
Submitted to *International Conference on Learning Representations (ICLR), 2025*

## Experience

**The University of North Carolina at Chapel Hill** — Remote
Research Intern — June 2023 – June 2024
Advisor: Prof. Tianlong Chen

**University of Science and Technology of China** — Hefei, China
Teaching Assistant — September 2022 – January 2023
Undergrad course - CS1001A Computer Programming A (C Language)

**University of Science and Technology of China** — Hefei, China
Undergrad Intern — July 2021
Advisor: Prof. Qi Liu, Prof. Enhong Chen

## Honors & Awards

1st Place of ACM/IEEE Quantum Computing for Drug Discovery Challenge November, 2023
Outstanding Graduates Scholarship, USTC June, 2023
Silver Medal in Kaggle Feedback Prize - Evaluating Student Writing March, 2022
Outstanding Student Scholarship, USTC November, 2020/21/22

## Services

**Reviewer**: NeurIPS (2024)
**Tutorial**: ICML 2024

## Skills

**Languages**: Mandarin (native), English (professional), German (junior)
**Programming Languages**: Python, C/C++, Bash
**Deep Learning Frameworks**: PyTorch, HuggingFace Transformers, DeepSpeed, Jax/Flax